

Tutor teaches. Learner moves on untested.

Job: confirm one concept immediately after a tutor explanation — before a misconception becomes expensive.

0 / 2,522

chatlog rows contained a real
proactive comprehension check


The widest gap is the missing check after teaching

Impact was compared before choosing the slice — not reverse-engineered after build.

Candidate	Reach	Observed signal	Decision
Comprehension check	369 / 369 users	98.4% of 1,261 turns teach; no follow-up check	CHOSEN
Summary/retrieval fix	102 / 369 users	51.1% failure for summary-like requests	Rejected
Missing citation fix	70 / 369 users	76% of direct answers lacked citations	Rejected

Cost of error: an unwanted quiz is skippable; an undetected misconception compounds until assessment.

A bounded state machine controls every release decision



STANDARD CASE

Day 1 · page 29 · top_p
Answer → one four-option quiz →
deterministic scoring

HARD CASE

Ask “2 + 2 = ?”
ABSTAINED · outside_course_scope

2

max generation attempts

Status clears the bar; semantic review stays open

Live run 7889852f · 28 preserved cases · OpenAI · immutable evidence files

LOCKED BAR: $\geq 80\%$ status + zero invention

**Groundedness · concept alignment · domain correctness =
HUMAN REVIEW REQUIRED**

96%

27 / 28 status matches

100%

18 / 18 valid schemas

G26

over-refused; expected OK

Human validation remains a submission gap

0

verified human feedback
responses recorded

HUMAN ACTION REQUIRED

- ✓ Test ≥ 5 people outside the team
- ✓ Include ≥ 2 willing users from CP1
- ✓ Record verbatim quotes with consent
- ✓ Link changes to feedback row IDs

Make semantic evidence as strong as schema evidence.

- 1 · Fix G26 over-refusal without regressing out-of-scope cases
- 2 · Complete two-reviewer grounding/concept/domain scoring
- 3 · Run ≥ 5 human sessions and ship only evidence-linked changes

**Prompts are not
proof.
Validation owns
release.**